

INT2 Group 14 Report

Group 14: Aous Abdulla, Ryan Doyle, Maddie Jones, Owen Lister, Sanjna Srinivasan

Abstract—Abstract — The task for this project was to construct and train a neural network which can classify images of flowers into one of 102 categories of species, given by the Oxford 102 Category Flower Dataset (Flowers-102). To complete the task, I used PyTorch to create a deep neural network that classifies the given flower images into the defined classes. This was achieved by first learning from tutorials via the PyTorch website, then adapting a basic network to my problem, and finally, fine tuning to improve accuracy. Overall, the network I designed achieved a classification accuracy of 69.9% on the flowers-102 test set.

I. INTRODUCTION

THE Flowers-102 dataset comprises images of flowers from 102 different species, with each species represented by at least 40 images. The task was to develop a deep neural network to classify these images into their respective categories.

Before Convolutional Neural Networks (CNNs), computer vision heavily relied on manual feature extraction methods, like edge detection, texture analysis, and colour histograms; whilst useful, these techniques had limitations in capturing complex patterns and variations within images.

In 1989, LeNet-1 emerged as one of the first CNN models for image classification [1]. It utilized backpropagation to recognize handwritten digits efficiently. This concept evolved with increased hidden and fully connected layers, leading to LeNet-5, which achieved lower error rates on test data, and finally to Boosted LeNet-4 in 1998, with even smaller error rates [2].

AlexNet, introduced in 2012, significantly advanced image recognition by incorporating data augmentation. Techniques such as image translations, horizontal reflections, and RGB intensity alterations, expanded the training set and reduced overfitting [3].

More recently, region-based CNNs (R-CNNs) have enhanced image classification by using image segmentation to draw boxes around different objects within an image, followed by CNNs to classify these Regions of Interest (ROI). Mask R-CNN is among the most effective models for image classification today [4].

In computer vision, image classification is essential alongside object detection, semantic segmentation, and image segmentation; posing a challenging problem due to the processing power required to analyse image details down to individual pixels. The Flowers-102 dataset presents additional complexities; for example, two flowers of the same species may vary in colour or flowering stage, yet their images need to be classified into the same category. To address this, I created and trained a CNN to classify images, applying various transformations, into their respective categories.

II. METHOD

I created a Convolutional Neural Network (CNN) – a specialised neural network for image processing and classification – to classify images from the Flowers-102 dataset. I began by loading the official training, validation and testing splits. Due to restrictions on pre-trained networks, I then design a basic CNN network. After experimenting with various architectures, I decided that data augmentation was necessary given the dataset's relatively small size (8000 images total, with 1020 in the training split).

For the training data, I applied the following transformations: Random Resized Crop: Images resized to 250x250 pixels, randomly cropped between 50% and 100%, then resized back to 250x250, Random Rotation: Images randomly rotated up to 55 degrees, Color Jitter: Adjusted brightness and contrast by up to 0.5, and saturation by up to 0.2, Random Horizontal and Vertical Flip: Probability of 0.5. Gaussian Blur: Kernel size of 5. Conversion to Tensor. For validation and test data, images were resized to 250x250 pixels and converted to tensors. No additional pre-processing steps were performed, and normalisation was omitted due to no observed benefit.

The CNN architecture includes six convolutional layers with increasing filters (16 to 256) and a kernel size of 3x3, with stride 1 and padding 1. Each layer used Batch Normalization to stabilize and accelerate training, PReLU activation for flexible negative slopes, Max Pooling to reduce spatial dimensions, and Dropout (0.1 probability) in the first, fifth, and sixth layers to prevent overfitting.

After the convolutional layers, a flattening layer converted the multi-dimensional output into a one-dimensional vector for the fully connected layers, which had dimensions: (256 * 3 * 3, 512), (512, 256), and (256, 102), finally outputting for the 102 classes.

The model was trained for up to 250 epochs, with early stopping if no improvement in validation accuracy was observed for 25 consecutive epochs. This approach balanced training time and accuracy. The training cycle to determine the ideal optimizer and learning rate scheduler, achieving the following recorded results on the test set:

Scheduler	Adam	AdamW	RMSprop
ReduceLROnPlateau	54.8%	42.3%	51.4%
CosineAnnealingLR	40.9%	46.6%	55.7%
LinearLR	56.5%	46.6%	53.5%

TABLE I
COMPARISON OF SCHEDULERS AND OPTIMIZERS

Following these performance metrics, the RMSprop optimizer with the CosineAnnealingLR scheduler was selected for its superior performance, and was used for all proceeding tests.

III. ARCHITECTURE

